

BEAST Lab

The Memory Hierarchy of Multi-core CPUs

May 5, 2022 | Josef Weidendorfer

Example: Intel Icelake



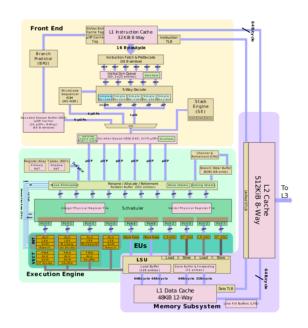
Two systems in BEAST

- 2 sockets Intel Xeon (Icelake) Platinum 8360Y
 - 2x 36 = 72 cores
 - 2x 512bit vector units per core (8 x DP FMA)
 - 2 threads per core ("Hyper-Threading")
 - 2.4 GHz base, Intel 10nm
- 512 GB main memory, 1.5 TB Optane NVRam

Links

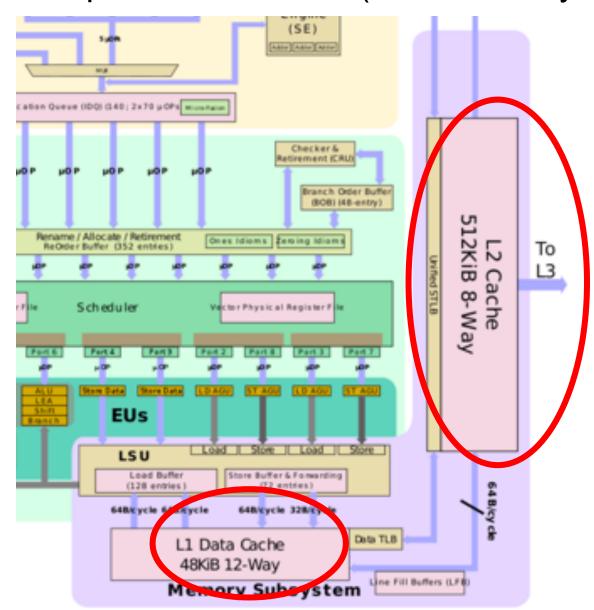
- https://en.wikichip.org/wiki/intel/microarchitectures/ice_lake_(server)
- https://en.wikichip.org/wiki/intel/microarchitectures/sunny_cove

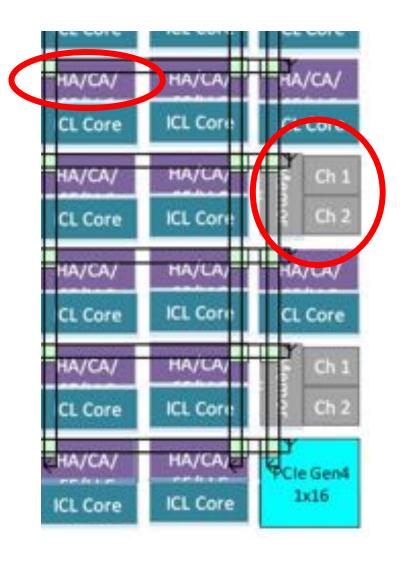




Example: Intel IceLake (similar to Skylake)

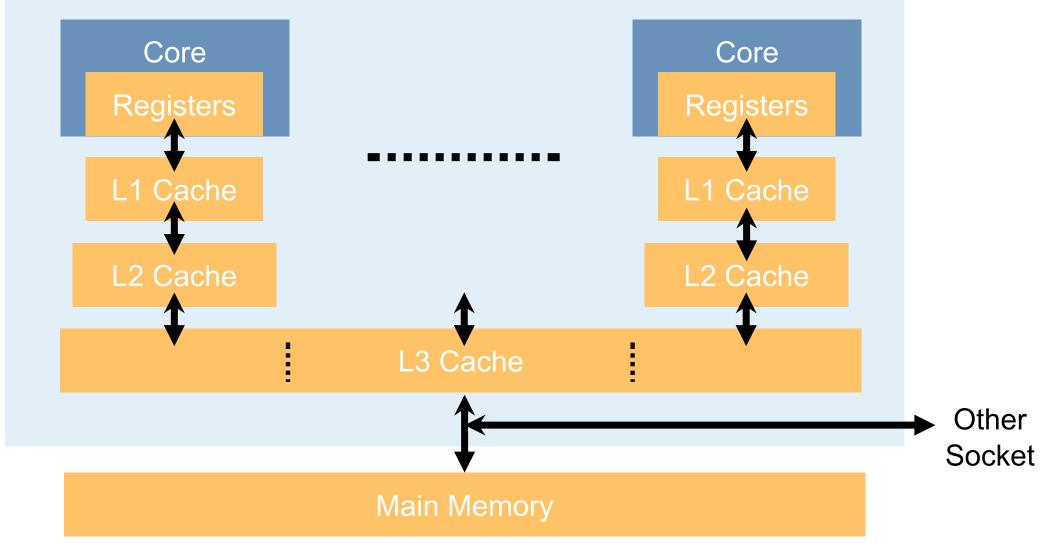






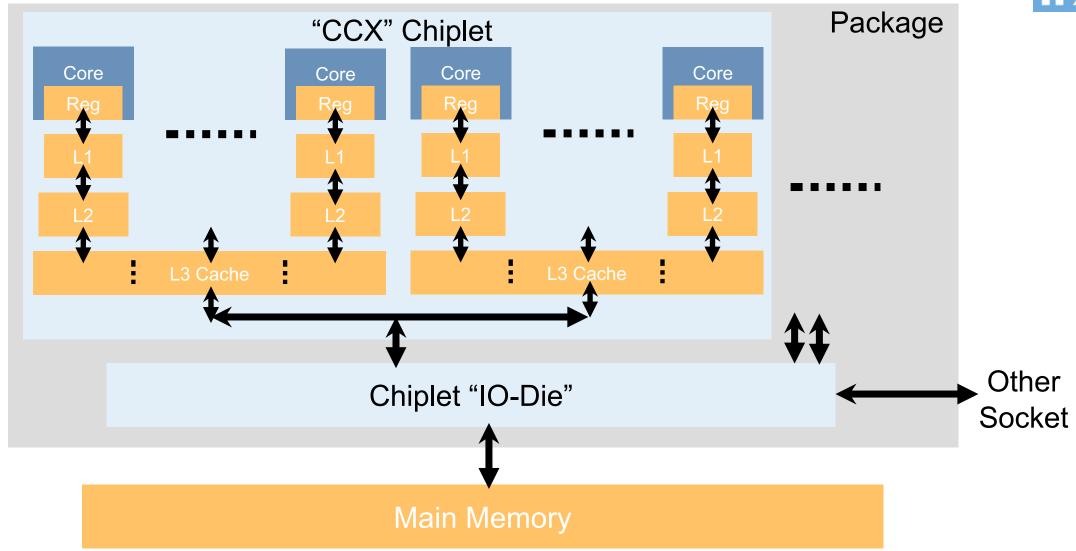
Example: Intel IceLake (similar to Skylake)





Example: AMD Rome





Cache and Cache-Levels: Motivation



Example: Laptop

- compute peak for 1 GHz CPU, 4 cores, 1 vector unit with 4 FP64 floats (FMA):
 10e9 /s * 4 * 4 * 2 Flops = 32 GFlop/s
- Triad (Assignment 1): A(i) = B(i) + C(i) * D(i)
 - per FP64 element: 3 * 8 B (Load) + 8 B (Store)
 - "Arithmetic Intensity": 2 Flops / 32 Bytes transferred = 0.067 Flop/B
 - required bandwidth to allow peak: 32 GFlop/s / 0.067 Flop/B = 512 GB/s
- Off-chip memory ~ 25 GB/s (Laptop) : 20x slower!
- Latency ~ 50ns: while accessing memory once, 1 core can do 50 * 8 = 400 Flops

Solution: keep recently accessed data on-chip in a "cache"

Cache and Cache-Levels



How to decide what data to keep in the cache?

Cache replacement policy: LRU (drop least-recently used if full)

Overhead: need to store both address and data in cache

- Programs often access data at near-side addresses → fetch data in blocks
- Can utilize parallel signal lanes to memory (multiple memory controllers + bursts)
- Implicit prefetch for near-side data
- Icelake / Rome / ThunderX2: 64 B, A64FX: 256 B

Performance of on-chip caches depend on cache size

- Smaller caches are faster → use a hierarchy of caches
- Icelake / Rome / ThunderX2: 3 levels, A64FX: 2 levels

Keep the Pipe filled: Memory Parallelism & Hardware Prefetching



Motivation: Assume synchronous, serialized memory accesses per core

- % of bandwidth usable?
- Best case: just memory accesses
 - 50 ns for 64 Bytes per core
 with sequential code (1 core): 64 / 50 ns ~ 1,2 GB/s (vs. 20 GB/s available on laptop)

Memory Parallelism

- caches support multiple outstanding loads (Icelake/Rome: > 8 per core)
- modern CPU cores can look ahead and trigger accesses out-of-order ("OoO")

Hardware Prefetching

- Cache predicts next memory accesses from recent access pattern
- Easier per core (L2), works very well for streams (up and down), can go wrong (!)

(Low-Level) Tuning Tips for Memory



Make use of caches

- Reorder your code such that you can reuse recently loaded data if possible Blocking / Tiling: instead of 10 times over full array, split into blocks and go 10x over each Q: How to take advantage of multiple cache levels (L1 / L2 / L3)?
- Choose a memory layout + access order such that all data in a cache line can be used before being evicted

Make use of hardware prefetcher

• For data too large to fit into cache: use stream accesses as much as possible

Make use memory parallelism

Independent accesses can be triggered in parallel with OoO

The Roofline Model – Visualization of CPU Limits



Given some characteristics of a given (parallel) code

- does it hit a hardware limit, or
- is there room for improvement?

Focus on limits

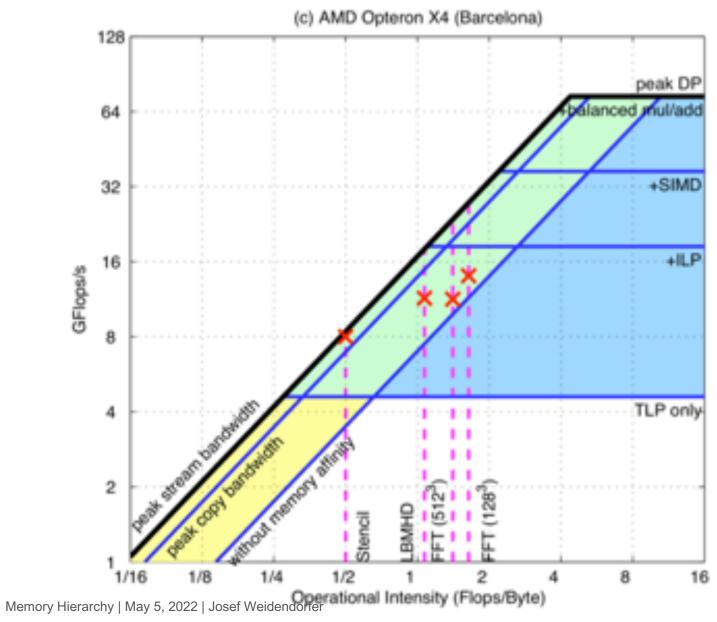
- memory bandwidth
- peak performance

Characteristic: Operational Intensity

- How many Flops are executed per Byte transferred to/from memory
- Can be different to Arithmetic Intensity (cache effects!)
 - Value is the same for vector triad for large vectors

The Roofline Model – Visualization of CPU Limits





[Williams et al: Roofline: An Insightful Visual Performance Model for Floating-Point Programs and Multicore Architectures]

The Roofline Model – Visualization of CPU Limits



Easy to use:

- Measure for your code: point in diagram
- How far are you from theoretical limits?

Extensions

- Limit vs. compute
 - L1 / L2 / L3 cache bandwidth, network, storage
 - Latency (L1 / L2 / L3 / main memory / network / storage)
 - Memory parallelism: how much accesses can be in-flight?
- Characterizations
 - Sensitivity to access latency
 - How many independent accesses could be triggered with X instruction look-ahead



Leibniz Supercomputing Centre of the Bavarian Academy of Sciences and Humanities